

```
#CIND820- Initial Code By Zion Yitbarek (500883564)
#Import libraries
import pandas as pd
import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.model_selection import cross_val_score, KFold
from sklearn.metrics import mean_absolute_error, r2_score

# Saving dataset to Google Colab
if 'df' in locals() or 'df' in globals():
    df.to_csv('retail_strategy_data.csv', index=False)
    print("File saved: retail_strategy_data.csv")
else:
    print("Error: DataFrame 'df' is not defined. Please ensure the cell defining 'df' (e.g., Cell WbCA3RZAllUL) is executed before this cell.")

File saved: retail_strategy_data.csv
```

```
# Inditex Group figures for CapEx Intensity and Gross Margin are assigned to both Zara and Massimo Dutti.
# Final Validated Dataset (observations: 2018-2023)
data = {
    'Brand': ['Zara', 'Massimo Dutti', 'Brunello Cucinelli'] * 6,
    'Year': sorted([2018] * 3 + [2019] * 3 + [2020] * 3 + [2021] * 3 + [2022] * 3 + [2023] * 3),
    'Revenue_M': [
        18021, 1802, 529.996, # 2018
        19564, 1900, 607.761, # 2019
        14129, 1197, 544.013, # 2020
        19586, 1653, 712.179, # 2021
        23761, 1593, 919.708, # 2022
        26050, 1839, 1139.42 # 2023
    ],
    #Core Financial Features (X)
    'Gross_Margin_Pct': [43.3, 43.3, 84.1, #2018
                         44.1, 44.1, 88.1, #2019
                         44.1, 44.1, 89.79, #2020
                         42.9, 42.9, 84, #2021
                         43, 43, 89.61, #2022
                         42.2, 42.2, 90], #2023
    'CapEx_M': [
        1609, 1609, 45, #2018
        1130, 1130, 52.6, #2019
        713, 713, 51.6, #2020
        1100, 1100, 61.6, #2021
        1400, 1400, 87.5, #2022
        1870, 1870, 79.1] #2023
}

df = pd.DataFrame(data)
df['Is_Luxury'] = df['Brand'].apply(lambda x: 1 if x == 'Brunello Cucinelli' else 0)

# Feature Engineering: Current Year's Growth (used as a feature)
df['YoY_Growth_Rate'] = df.groupby('Brand')['Revenue_M'].pct_change() * 100

# Target Variable (Y): Next Year's Growth (Measure of Resilience)
df['Target_Next_Year_Growth'] = df.groupby('Brand')['YoY_Growth_Rate'].shift(-1)
```

```
# Stating the three required machine learning algorithms I am using.
models = {
    "Linear Regression (Baseline)": LinearRegression(),
    "Random Forest Regressor": RandomForestRegressor(n_estimators=100, random_state=42),
    "Gradient Boosting Regressor": GradientBoostingRegressor(n_estimators=100, learning_rate=0.1, random_state=42)
}
# Defined the Cross-Validation
cv = KFold(n_splits=4, shuffle=True, random_state=42)
results = []
# Looping through models to observe performance (MAE and R2)
for name, model in models.items():
    # Mean Absolute Error (MAE) and R-squared (R2)
    mae_scores = cross_val_score(model, X, Y, cv=cv, scoring='neg_mean_absolute_error')
    r2_scores = cross_val_score(model, X, Y, cv=cv, scoring='r2')
```

```

mean_mae = -mae_scores.mean()

mean_r2 = r2_scores.mean()
results.append({
    'Algorithm': name,
    'Cross-Validated MAE': f"{mean_mae:.2f}",
    'Cross-Validated R2': f"{mean_r2:.2f}"
})
# Display the results
results_df = pd.DataFrame(results)
print("--- Table 1: Comparative Cross-Validation Results ---")
print(results_df.to_markdown(index=False))
rf_model = models["Random Forest Regressor"]
rf_model.fit(X, Y)
feature_importances = pd.Series(rf_model.feature_importances_, index=X.columns).sort_values(ascending=False)
print("\n--- Table 2: Random Forest Feature Importance (Insights) ---")
print(feature_importances.to_markdown(numalign="left", stralign="left"))

```

--- Table 1: Comparative Cross-Validation Results ---		
Algorithm	Cross-Validated MAE	Cross-Validated R2
Linear Regression (Baseline)	71.91	-33.37
Random Forest Regressor	19.56	-0.32
Gradient Boosting Regressor	22.11	-1.67

--- Table 2: Random Forest Feature Importance (Insights) ---	
0	
CapEx_M	0.373721
YoY_Growth_Rate	0.319491
Revenue_M	0.15579
GrossMargin_Pct	0.141286
Is_Luxury	0.0097123

```

correlation_matrix = df.corr(numeric_only=True)
print("--- Correlation Matrix ---")
print(correlation_matrix.round(3).to_markdown())

```

--- Correlation Matrix ---								
	Year	Revenue_M	GrossMargin_Pct	CapEx_M	Is_Luxury	YoY_Growth_Rate	Target_N	
Year	1	0.109	0.006	0.13	-0	0.354		
Revenue_M	0.109	1	-0.522	0.526	-0.52	0.07		
GrossMargin_Pct	0.006	-0.522	1	-0.892	0.997	0.208		
CapEx_M	0.13	0.526	-0.892	1	-0.886	-0.006		
Is_Luxury	-0	-0.52	0.997	-0.886	1	0.234		
YoY_Growth_Rate	0.354	0.07	0.208	-0.006	0.234	1		
Target_Next_Year_Growth	0.354	-0.088	0.23	-0.332	0.234	-0.307		

```

growth_summary = df.dropna(subset=['YoY_Growth_Rate']).groupby('Brand')['YoY_Growth_Rate'].agg(['mean', 'std'])

```

```

print("\n--- Comparative Growth Summary (2019-2023) ---")
print(growth_summary.round(2).to_markdown())

```

```

import matplotlib.pyplot as plt

```

```

df_plot = df.dropna(subset=['YoY_Growth_Rate'])

```

```

plt.figure(figsize=(10, 6))
plt.title('Comparative YoY Revenue Growth Rate (2019-2023)')
plt.xlabel('Year')
plt.ylabel('YoY Growth Rate (%)')
plt.grid(axis='y', linestyle='--')

for brand in df_plot['Brand'].unique():
    brand_data = df_plot[df_plot['Brand'] == brand]
    plt.plot(brand_data['Year'], brand_data['YoY_Growth_Rate'],
             marker='o', label=brand)

```

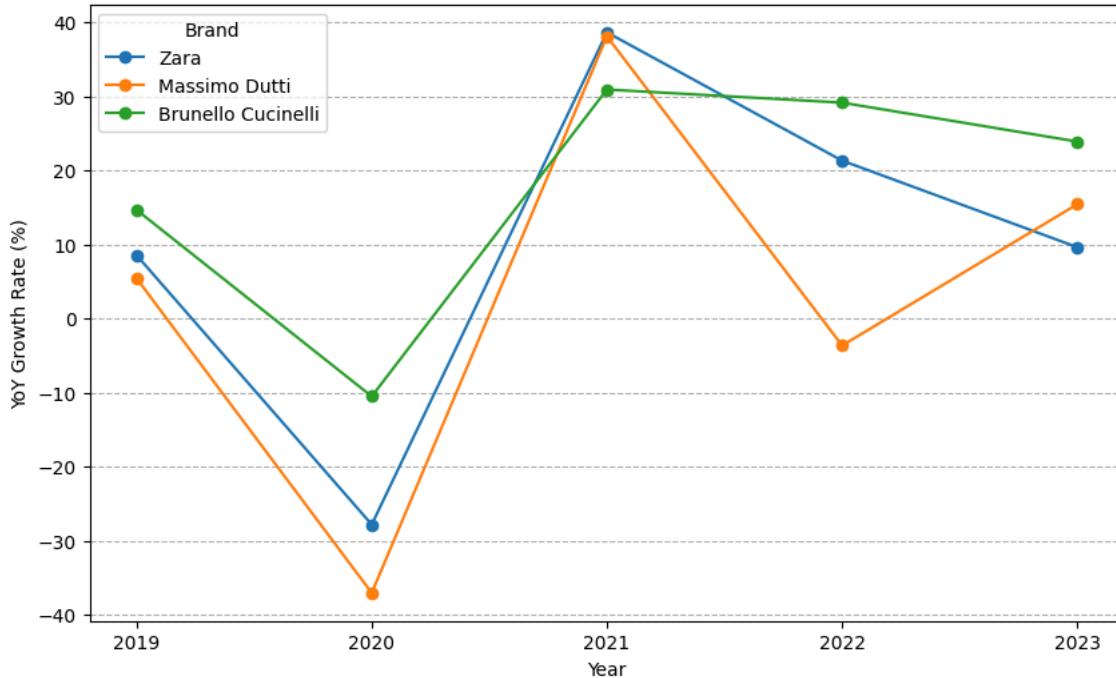
```

plt.legend(title='Brand')
plt.xticks(df_plot['Year'].unique())
plt.show()

```

--- Comparative Growth Summary (2019-2023) ---		
Brand	mean	std
Brunello Cucinelli	17.63	16.94
Massimo Dutti	3.67	27.54
Zara	10.07	24.38

Comparative YoY Revenue Growth Rate (2019-2023)



```

import numpy as np
from sklearn.cluster import KMeans

clustering_features = ['Gross_Margin_Pct', 'CapEx_M', 'YoY_Growth_Rate']
X_cluster = df[clustering_features].dropna()

print("Data prepared for K-Means. Total observations to cluster:", X_cluster.shape[0])

from sklearn.cluster import KMeans

# Define the number of clusters (k=3)
k = 3

# 2. Run K-Means
# X_cluster is the clean data
kmeans = KMeans(n_clusters=k, random_state=42, n_init='auto')
cluster_labels = kmeans.fit_predict(X_cluster)

# add cluster labels back to the DataFrame
df.loc[X_cluster.index, 'Cluster'] = cluster_labels.astype(int)

# Calculate mean values of features for each cluster
clustering_features = ['Gross_Margin_Pct', 'CapEx_M', 'YoY_Growth_Rate']
cluster_analysis = df.groupby('Cluster')[clustering_features + ['Brand']].agg({
    'Gross_Margin_Pct': 'mean',
    'CapEx_M': 'mean',
    'YoY_Growth_Rate': 'mean',
    'Brand': 'first' # Use 'first' to get a representative brand for the output table
})

# Add count of observations per cluster
cluster_analysis['Count'] = df.groupby('Cluster').size()
cluster_analysis = cluster_analysis.rename(columns={'Brand': 'Representative_Brand'})

print("--- K-Means Cluster Analysis (k=3) ---")
print(cluster_analysis.round(2).to_markdown())

# show what brands are in what clusters

```

```
full_cluster_composition = df.dropna(subset=['Cluster']).groupby('Cluster')['Brand'].value_counts().unstack(fill_value=0)
print("\n--- Cluster Composition by Brand ---")
print(full_cluster_composition.to_markdown())
```

Data prepared for K-Means. Total observations to cluster: 15

--- K-Means Cluster Analysis (k=3) ---

Cluster	Gross_Margin_Pct	CapEx_M	YoY_Growth_Rate	Representative_Brand	Count
0	88.3	66.48	17.63	Brunello Cucinelli	5
1	42.2	1870	12.54	Zara	2
2	43.52	1085.75	5.45	Zara	8

--- Cluster Composition by Brand ---

Cluster	Brunello Cucinelli	Massimo Dutti	Zara
0	5	0	0
1	0	1	1
2	0	4	4

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
# Inditex Group figures for CapEx Intensity and Gross Margin are assigned to both Zara and Massimo Dutti.
```

```
# Final Validated Dataset (observations: 2018-2023)
```

```
data = {
```

```
    'Brand': ['Zara', 'Massimo Dutti', 'Brunello Cucinelli'] * 6,
```

```
    'Year': sorted([2018] * 3 + [2019] * 3 + [2020] * 3 + [2021] * 3 + [2022] * 3 + [2023] * 3),
```

```
'Revenue_M': [
```

```
    18021, 1802, 529.996, # 2018
```

```
    19564, 1900, 607.761, # 2019
```

```
    14129, 1197, 544.013, # 2020
```

```
    19586, 1653, 712.179, # 2021
```

```
    23761, 1593, 919.708, # 2022
```

```
    26050, 1839, 1139.42 # 2023
```

```
],
```

```
#Core Financial Features (X)
```

```
'Gross_Margin_Pct': [43.3, 43.3, 84.1, #2018
```

```
    44.1, 44.1, 88.1, #2019
```

```
    44.1, 44.1, 89.79, #2020
```

```
    42.9, 42.9, 84, #2021
```

```
    43, 43, 89.61, #2022
```

```
    42.2, 42.2, 90], #2023
```

```
}
```

```
df = pd.DataFrame(data)
```

```
df['Is_Luxury'] = df['Brand'].apply(lambda x: 1 if x == 'Brunello Cucinelli' else 0)
```

```
# Feature Engineering: Current Year's Growth (used as a feature)
```

```
df['YoY_Growth_Rate'] = df.groupby('Brand')['Revenue_M'].pct_change() * 100
```

```
# Target Variable (Y): Next Year's Growth (Measure of Resilience)
```

```
df['Target_Next_Year_Growth'] = df.groupby('Brand')['YoY_Growth_Rate'].shift(-1)
```

```
# Aggregate Zara and Massimo Dutti into 'Inditex'
```

```
inditex_capex = df[df['Brand'].isin(['Zara', 'Massimo Dutti'])].groupby('Year')['CapEx_M'].sum().reset_index()
inditex_capex['Brand'] = 'Inditex'
```

```
# Get data for Brunello Cucinelli
```

```
brune_capex = df[df['Brand'] == 'Brunello Cucinelli'][['Year', 'CapEx_M', 'Brand']]
```

```
# Combine the aggregated Inditex data and Brunello Cucinelli data
plot_df = pd.concat([inditex_capex, brune_capex], ignore_index=True)
```

```
# Set plot style
```

```
sns.set_style("whitegrid")
```

```
plt.figure(figsize=(12, 6))
```

```
# Plot CapEx_M for each brand over the years from the new plot_df
for brand in plot_df['Brand'].unique():
    ...
```

```
brand_data = plot_dt[plot_dt['Brand'] == brand]
plt.plot(brand_data['Year'], brand_data['CapEx_M'], marker='o', label=brand)

plt.title('Capital Expenditure (CapEx_M) by Brand Over Years (Inditex Combined)')
plt.xlabel('Year')
plt.ylabel('CapEx (in Millions)')
plt.xticks(df['Year'].unique())
plt.legend(title='Brand')
plt.grid(True)
plt.tight_layout()
plt.show()
```

